

# Capstone Project: Instant Health Alert System - Final-Project Submission

## EMR instances

- EMR cluster (with Spark, Hive, Sqoop,Hbase)

Cluster info

Cluster ID  
j-3H2EM4F31JT0Z

Cluster configuration  
Instance groups

Capacity  
1 Primary 0 Core 0 Task

Applications

Amazon EMR version  
emr-5.30.1

Installed applications  
HBase 1.4.13, Hadoop 2.8.5, Hive 2.3.6, Hue 4.6.0, JupyterHub 1.1.0, Livy 0.7.0, Spark 2.4.5, Sqoop 1.4.7, Zeppelin 0.8.2

Cluster management

Log destination in Amazon S3  
aws-logs-284159339388-us-east-1/elasticmapreduce

Persistent application UIs  
Spark History Server  
YARN timeline server  
Tez UI

Primary node public DNS  
ec2-44-201-201-67.compute-1.amazonaws.com  
Connect to the Primary node using SSH  
Connect to the Primary node using SSM

Status and time

Status  
Waiting

Creation time  
November 11, 2024, 12:37 (UTC+05:30)

Elapsed time  
1 hour, 22 minutes

## EMR Hardware Configuration (with 1 master node)

Instance groups (1) info

With the instance groups configuration, each node type consists of the same instance type and the same purchasing option for instances: On-Demand or Spot.

Find instances by status

Find resources by ID or type; or search for text within loaded results

Terminate instance

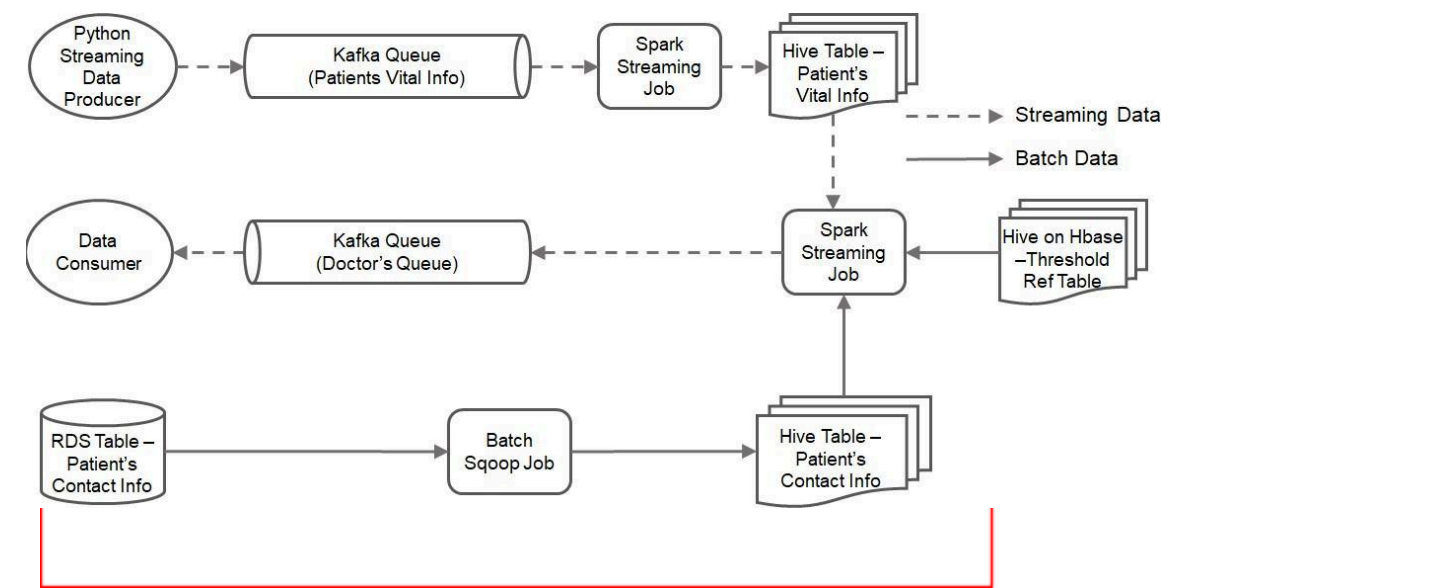
Resize instance group

Add task instance group

Type and name	ID	Status	Instances	Purchasing option and p...	EBS size (GiB)	EC2 Instance ID	Public DNS name	Public IP address
Primary	ig-2G43Z5CP3OT53	Running	1	On-Demand	-	-	-	-
m4.xlarge	ci-03796753D8XUE4Z...	Running	-	On-Demand (-)	40 (1 volume)	i-0a9199c4e15b52651	ec2-44-201-201-67.cor	44.201.20.67

## PART 1:

We will import patient contact information from the RDS using Sqoop and load the data into a Hive table for further processing.



1. Import Patient Contact Info records to HDFS using Sqoop
2. Creating an external Hive table for storing Patient's Contact Info

## SQOOP SETUP

Following steps are followed to setup Sqoop on EMR Cluster

1. To install the MySQL connector jar file.

```
wget https://de-mysql-connector.s3.amazonaws.com/mysql-connector-java-8.0.25.tar.gz
```

```
[hadoop@ip-172-31-12-140 ~]$ wget https://de-mysql-connector.s3.amazonaws.com/mysql-connector-java-8.0.25.tar.gz
--2024-11-11 07:27:39-- https://de-mysql-connector.s3.amazonaws.com/mysql-connector-java-8.0.25.tar.gz
Resolving de-mysql-connector.s3.amazonaws.com (de-mysql-connector.s3.amazonaws.com)... 52.217.199.41, 52.217.195.17, 3.5.23.166, ...
Connecting to de-mysql-connector.s3.amazonaws.com (de-mysql-connector.s3.amazonaws.com)|52.217.199.41|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 4079310 (3.9M) [application/x-gzip]
Saving to: 'mysql-connector-java-8.0.25.tar.gz'

100%[=====] 4,079,310 --.-K/s in 0.06s

2024-11-11 07:27:39 (60.0 MB/s) - 'mysql-connector-java-8.0.25.tar.gz' saved [4079310/4079310]
```

2. Extract the MySQL connector tar file

```
tar -xvf mysql-connector-java-8.0.25.tar.gz
```

```
[hadoop@ip-172-31-12-140 ~]$ wget https://de-mysql-connector.s3.amazonaws.com/mysql-connector-java-8.0.25.tar.gz
--2024-11-11 07:27:39-- https://de-mysql-connector.s3.amazonaws.com/mysql-connector-java-8.0.25.tar.gz
Resolving de-mysql-connector.s3.amazonaws.com (de-mysql-connector.s3.amazonaws.com)... 52.217.199.41, 52.217.195.17, 3.5.23.166, ...
Connecting to de-mysql-connector.s3.amazonaws.com (de-mysql-connector.s3.amazonaws.com)|52.217.199.41|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 4079310 (3.9M) [application/x-gzip]
Saving to: 'mysql-connector-java-8.0.25.tar.gz'

100%[=====] 4,079,310 --.-K/s in 0.06s

2024-11-11 07:27:39 (60.0 MB/s) - 'mysql-connector-java-8.0.25.tar.gz' saved [4079310/4079310]

[hadoop@ip-172-31-12-140 ~]$ tar -xvf mysql-connector-java-8.0.25.tar.gz
mysql-connector-java-8.0.25/
mysql-connector-java-8.0.25/src/
mysql-connector-java-8.0.25/src/build/
mysql-connector-java-8.0.25/src/build/java/
mysql-connector-java-8.0.25/src/build/java/documentation/
mysql-connector-java-8.0.25/src/build/java/instrumentation/
mysql-connector-java-8.0.25/src/build/misc/
mysql-connector-java-8.0.25/src/build/misc/debian.in/
mysql-connector-java-8.0.25/src/build/misc/debian.in/source/
mysql-connector-java-8.0.25/src/demo/
mysql-connector-java-8.0.25/src/demo/java/
mysql-connector-java-8.0.25/src/demo/java/demo/
mysql-connector-java-8.0.25/src/demo/java/demo/x/
mysql-connector-java-8.0.25/src/demo/java/demo/x/devapi/
mysql-connector-java-8.0.25/src/generated/
mysql-connector-java-8.0.25/src/generated/java/
mysql-connector-java-8.0.25/src/generated/java/com/
mysql-connector-java-8.0.25/src/generated/java/com/mysql/
mysql-connector-java-8.0.25/src/generated/java/com/mysql/cj/
mysql-connector-java-8.0.25/src/generated/java/com/mysql/cj/x/
mysql-connector-java-8.0.25/src/generated/java/com/mysql/cj/x/protobuf/
mysql-connector-java-8.0.25/src/legacy/
mysql-connector-java-8.0.25/src/legacy/java/
mysql-connector-java-8.0.25/src/legacy/java/com/
```

3. Navigate to the MySQL Connector directory created in the previous step, and copy it to the Sqoop library to complete the installation.

```
cd mysql-connector-java-8.0.25/
sudo cp mysql-connector-java-8.0.25.jar /usr/lib/sqoop/lib/
```

```
[hadoop@ip-172-31-12-140 ~]$ cd mysql-connector-java-8.0.25/
[hadoop@ip-172-31-12-140 mysql-connector-java-8.0.25]$ sudo cp mysql-connector-java-8.0.25.jar /usr/lib/sqoop/lib/
[hadoop@ip-172-31-12-140 mysql-connector-java-8.0.25]$ mysql config installation
```

#### 4. Set up MySQL on your EMR cluster (Inside this folder mysql-connector-java-8.0.25)

mysql\_secure\_installation

Enter current password for root (enter for none): ENTER

Set root password [Y/n] Y New password: 123

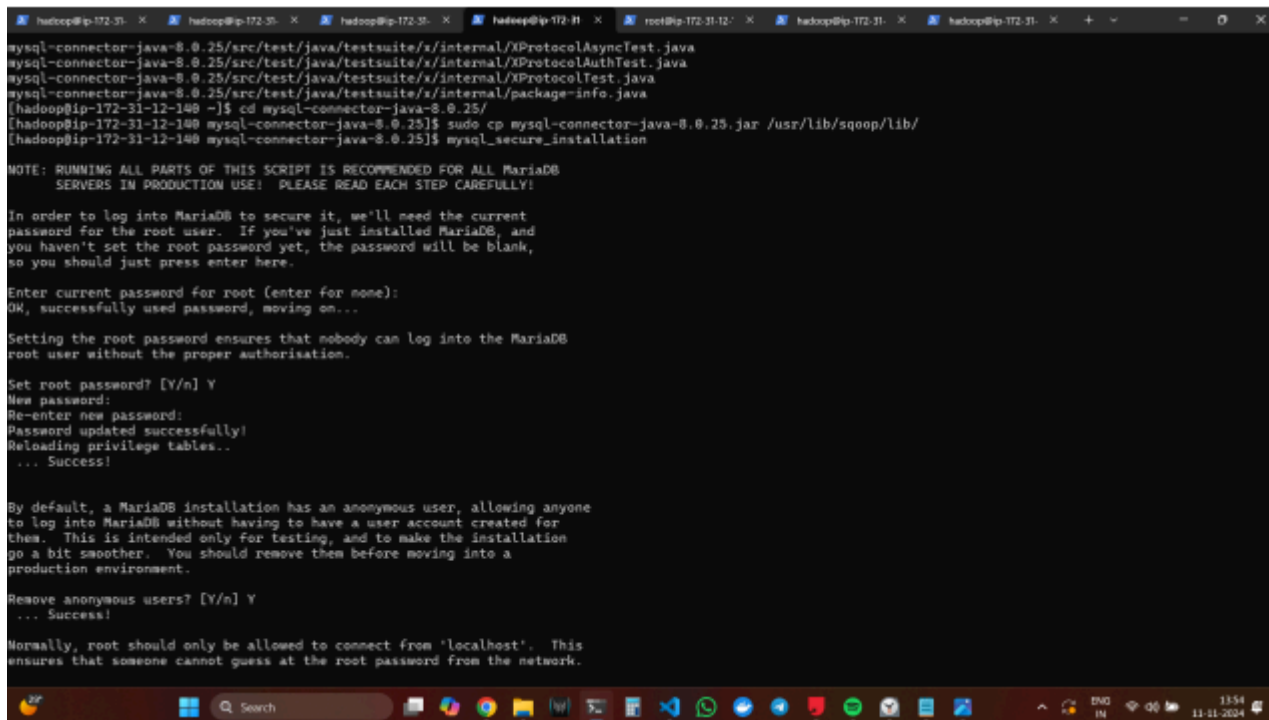
Re-enter password: 123

Remove anonymous users [Y/n] Y

Disallow root login remotely [Y/n] n

Remove test database and access to it [Y/n] Y

Reload privilege tables now [Y/n] Y



```
mysql-connector-java-8.0.25/src/test/java/testsuite/x/internal/XProtocolAsyncTest.java
mysql-connector-java-8.0.25/src/test/java/testsuite/x/internal/XProtocolAuthTest.java
mysql-connector-java-8.0.25/src/test/java/testsuite/x/internal/XProtocolTest.java
mysql-connector-java-8.0.25/src/test/java/testsuite/x/internal/package-info.java
[hadoop@ip-172-31-12-140 ~]$ cd mysql-connector-java-8.0.25/
[hadoop@ip-172-31-12-140 mysql-connector-java-8.0.25]$ sudo cp mysql-connector-java-8.0.25.jar /usr/lib/sqoop/lib/
[hadoop@ip-172-31-12-140 mysql-connector-java-8.0.25]$ mysql_secure_installation

NOTE: RUNNING ALL PARTS OF THIS SCRIPT IS RECOMMENDED FOR ALL MariaDB
SERVERS IN PRODUCTION USE! PLEASE READ EACH STEP CAREFULLY!

In order to log into MariaDB to secure it, we'll need the current
password for the root user. If you've just installed MariaDB, and
you haven't set the root password yet, the password will be blank,
so you should just press enter here.

Enter current password for root (enter for none):
OK, successfully used password, moving on...

Setting the root password ensures that nobody can log into the MariaDB
root user without the proper authorisation.

Set root password? [Y/n] Y
New password:
Re-enter new password:
Password updated successfully!
Reloading privilege tables..
... Success!

By default, a MariaDB installation has an anonymous user, allowing anyone
to log into MariaDB without having to have a user account created for
them. This is intended only for testing, and to make the installation
go a bit smoother. You should remove them before moving into a
production environment.

Remove anonymous users? [Y/n] Y
... Success!

Normally, root should only be allowed to connect from 'localhost'. This
ensures that someone cannot guess at the root password from the network.

Reload privilege tables now? [Y/n] Y
... Success!

Cleaning up...

All done! If you've completed all of the above steps, your MariaDB
installation should now be secure.

Thanks for using MariaDB!
```

By default, MariaDB comes with a database named 'test' that anyone can access. This is also intended only for testing, and should be removed before moving into a production environment.

Remove test database and access to it? [Y/n] Y

- Dropping test database...

... Success!

- Removing privileges on test database...

... Success!

Reloading the privilege tables will ensure that all changes made so far will take effect immediately.

Reload privilege tables now? [Y/n] Y

... Success!

Cleaning up...

All done! If you've completed all of the above steps, your MariaDB installation should now be secure.

Thanks for using MariaDB!

6. After this a prompt for password will be shown hit 123 and grant root privileges to the other user and restart mariadb

```
mysql -u root -p
```

```
[hadoop@ip-172-31-83-130 mysql-connector-java-8.0.25]$ mysql -u root -p
Enter password:
Welcome to the MariaDB monitor.  Commands end with ; or \g.
Your MariaDB connection id is 66
Server version: 5.5.68-MariaDB MariaDB Server

Copyright (c) 2000, 2018, Oracle, MariaDB Corporation Ab and others.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

MariaDB [(none)]> GRANT ALL PRIVILEGES ON *.* TO 'root'@'%' identified by '123'
```

7. Inside MariaDB (MariaDB > )

Following queries need to be run for granting all privileges to the root user.

```
GRANT ALL PRIVILEGES ON *.* TO 'root'@'%' identified by '123' WITH GRANT OPTION;
flush privileges;
exit;
```

```
MariaDB [(none)]> GRANT ALL PRIVILEGES ON *.* TO 'root'@'%' identified by '123'
WITH GRANT OPTION;
Query OK, 0 rows affected (0.00 sec)

MariaDB [(none)]> flush privileges;
Query OK, 0 rows affected (0.00 sec)

MariaDB [(none)]> exit;
Bye
```

8. Restart the MySQL service to finish setting up MySQL. (Inside this folder mysql-connector-java-8.0.25 )

```
sudo service mariadb restart
```

```
[hadoop@ip-172-31-83-130 mysql-connector-java-8.0.25]$ sudo service mariadb rest
art
Redirecting to /bin/systemctl restart mariadb.service
```

## [Sqoop Commands](#)

1. Import data to HDFS

```
sqoop import --connect
jdbc:mysql://upgraddetest.cyaieic9bmnf.us-east-1.rds.amazonaws.com/testdatabase --table
patients_information --username student --password STUDENT123 --target-dir
/user/hadoop/health-alert/patients-contact-info -m 1
```

```
[hadoop@ip-172-31-12-140 ~]$ sqoop import --connect jdbc:mysql://upgradetest.cyaie1c9bmhf.us-east-1.rds.amazonaws.com/testdatabase --table patients_informa
tion --username student --password STUDENT123 --target-dir /user/hadoop/health-alert/patients-contact-info --m 1
Warning: /usr/lib/sqoop/.accumulo does not exist! Accumulo imports will fail.
Please set $ACCUMULO_HOME to the root of your Accumulo installation.
24/11/11 07:29:16 INFO sqoop.Sqoop: Running Sqoop version: 1.4.7
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/usr/lib/hadoop/lib/slf4j-log4j12-1.7.10.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/share/aws/redshift/jdbc/redshift-jdbc42-1.2.37.1861.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/lib/hive/lib/log4j-slf4j-impl-2.6.2.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
24/11/11 07:29:16 WARN tool.BaseSqoopTool: Setting your password on the command-line is insecure. Consider using -P instead.
24/11/11 07:29:16 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.
24/11/11 07:29:16 INFO tool.CodeGenTool: Beginning code generation
Loading class 'com.mysql.jdbc.Driver'. This is deprecated. The new driver class is 'com.mysql.cj.jdbc.Driver'. The driver is automatically registered via th
e SPI and manual loading of the driver class is generally unnecessary.
24/11/11 07:29:17 INFO manager.MySQLManager: Executing SQL statement: SELECT t.* FROM 'patients_information' AS t LIMIT 1
24/11/11 07:29:17 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM 'patients_information' AS t LIMIT 1
24/11/11 07:29:17 INFO orm.CompilationManager: HADOOP_MAPRED_HOME is /usr/lib/hadoop-mapreduce
Note: /tmp/sqoop-hadoop/compile/adel1f48bdd66c572a48f8b315c16c818/patients_information.java uses or overrides a deprecated API.
Note: Recompile with -Xlint:deprecation for details.
24/11/11 07:29:20 INFO orm.CompilationManager: Writing jar file: /tmp/sqoop-hadoop/compile/adel1f48bdd66c572a48f8b315c16c818/patients_information.jar
24/11/11 07:29:20 WARN manager.MySQLManager: It looks like you are importing from mysql.
24/11/11 07:29:20 WARN manager.MySQLManager: This transfer can be faster! Use the --direct
24/11/11 07:29:20 WARN manager.MySQLManager: option to exercise a MySQL-specific fast path.
24/11/11 07:29:20 INFO manager.MySQLManager: Setting zero DATETIME behavior to convertToNull (mysql)
24/11/11 07:29:20 INFO mapreduce.ImportJobBase: Beginning import of patients_information
24/11/11 07:29:21 INFO Configuration.deprecation: mapred.jar is deprecated. Instead, use mapreduce.job.jar
24/11/11 07:29:21 INFO Configuration.deprecation: mapred.map.tasks is deprecated. Instead, use mapreduce.job.maps
24/11/11 07:29:21 INFO client.RMProxy: Connecting to ResourceManager at ip-172-31-12-140.ec2.internal/172.31.12.140:8032
24/11/11 07:29:25 INFO db.DBInputFormat: Using read committed transaction isolation
24/11/11 07:29:25 INFO mapreduce.JobSubmitter: number of splits:1
24/11/11 07:29:25 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1731309572038_0001
24/11/11 07:29:25 INFO impl.YarnClientImpl: Submitted application application_1731309572038_0001
24/11/11 07:29:26 INFO mapreduce.Job: The url to track the job: http://ip-172-31-12-140.ec2.internal:20888/proxy/application_1731309572038_0001/
24/11/11 07:29:26 INFO mapreduce.Job: Running job: job_1731309572038_0001
24/11/11 07:29:35 INFO mapreduce.Job: Job job_1731309572038_0001 running in uber mode : false
24/11/11 07:29:35 INFO mapreduce.Job: map 0% reduce 0%
```

```
24/11/11 07:29:35 INFO mapreduce.Job: Job job_1731309572038_0001 running in uber mode : false
24/11/11 07:29:35 INFO mapreduce.Job: map 0% reduce 0%
24/11/11 07:29:43 INFO mapreduce.Job: map 100% reduce 0%
24/11/11 07:29:44 INFO mapreduce.Job: Job job_1731309572038_0001 completed successfully
24/11/11 07:29:44 INFO mapreduce.Job: Counters: 30
File System Counters
  FILE: Number of bytes read=0
  FILE: Number of bytes written=189970
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=87
  HDFS: Number of bytes written=230
  HDFS: Number of read operations=4
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=2
Job Counters
  Launched map tasks=1
  Other local map tasks=1
  Total time spent by all maps in occupied slots (ms)=282432
  Total time spent by all reduces in occupied slots (ms)=0
  Total time spent by all map tasks (ms)=5884
  Total vcore-milliseconds taken by all map tasks=5884
  Total megabyte-milliseconds taken by all map tasks=9037824
Map-Reduce Framework
  Map input records=5
  Map output records=5
  Input split bytes=87
  Spilled Records=0
  Failed Shuffles=0
  Merged Map outputs=0
  GC time elapsed (ms)=86
  CPU time spent (ms)=3350
  Physical memory (bytes) snapshot=358223872
  Virtual memory (bytes) snapshot=3298961456
  Total committed heap usage (bytes)=326107136
File Input Format Counters
  Bytes Read=0
File Output Format Counters
  Bytes Written=230
24/11/11 07:29:44 INFO mapreduce.ImportJobBase: Transferred 230 bytes in 22.889 seconds (10.0485 bytes/sec)
```

```
Other local map tasks=1
Total time spent by all maps in occupied slots (ms)=161328
Total time spent by all reduces in occupied slots (ms)=0
Total time spent by all map tasks (ms)=3361
Total vcore-milliseconds taken by all map tasks=3361
Total megabyte-milliseconds taken by all map tasks=5162496
Map-Reduce Framework
  Map input records=5
  Map output records=5
  Input split bytes=87
  Spilled Records=0
  Failed Shuffles=0
  Merged Map outputs=0
  GC time elapsed (ms)=67
  CPU time spent (ms)=1890
  Physical memory (bytes) snapshot=261730304
  Virtual memory (bytes) snapshot=3281002496
  Total committed heap usage (bytes)=247463936
File Input Format Counters
  Bytes Read=0
File Output Format Counters
  Bytes Written=230
23/03/25 07:11:39 INFO mapreduce.ImportJobBase: Transferred 230 bytes in 20.9571
seconds (10.9748 bytes/sec)
23/03/25 07:11:39 INFO mapreduce.ImportJobBase: Retrieved 5 records.
```

## Hive table creation (for Patients\_Contact\_Info)

- Open Hive shell.
- Create a database *health*

```
create database health;
```

- Use database patient\_health\_care

```
use health;
```

```
Logging initialized using configuration in file:/etc/hive/conf.dist/hive-log4j2.properties Async: false
hive> create database health;
OK
Time taken: 0.48 seconds
hive> use health;
OK
Time taken: 0.02 seconds
hive> |
```

- Create external table named *Patients\_Contact\_Info*

```
CREATE EXTERNAL TABLE IF NOT EXISTS health.patients_contact_info (
  patientid int,
  patientname string,
  patientaddress string,
  phone_number string,
  admitted_ward int,
  age int,
  other_details string
)
row format delimited
fields terminated by ','
lines terminated by '\n'
location '/user/hadoop/health-alert/patients-contact-info';
```

```
Input split bytes=87
Spilled Records=0
Failed Shuffles=0
Merged Map outputs=0
GC time elapsed (ms)=86
CPU time spent (ms)=3350
Physical memory (bytes) snapshot=358223872
Virtual memory (bytes) snapshot=3298963456
Total committed heap usage (bytes)=326107136

File Input Format Counters
  Bytes Read=0
File Output Format Counters
  Bytes Written=230
24/11/11 07:29:44 INFO mapreduce.ImportJobBase: Transferred 230 bytes in 22.889 seconds (10.0485 bytes/sec)
24/11/11 07:29:44 INFO mapreduce.ImportJobBase: Retrieved 5 records.
[hadoop@ip-172-31-12-140 ~]$ hive

Logging initialized using configuration in file:/etc/hive/conf.dist/hive-log4j2.properties Async: false
hive> create database health;
OK
Time taken: 1.07 seconds
hive> use health;
OK
Time taken: 0.057 seconds
hive> CREATE EXTERNAL TABLE IF NOT EXISTS health.patients_contact_info (
  > patientid int,
  > patientname string,
  > patientaddress string,
  > phone_number string,
  > admitted_ward int,
  > age int,
  > other_details string
  > )
  > row format delimited
  > fields terminated by ','
  > lines terminated by '\n'
  > location '/user/hadoop/health-alert/patients-contact-info';
OK
Time taken: 0.297 seconds
hive> select * from patients_contact_info;
OK
```



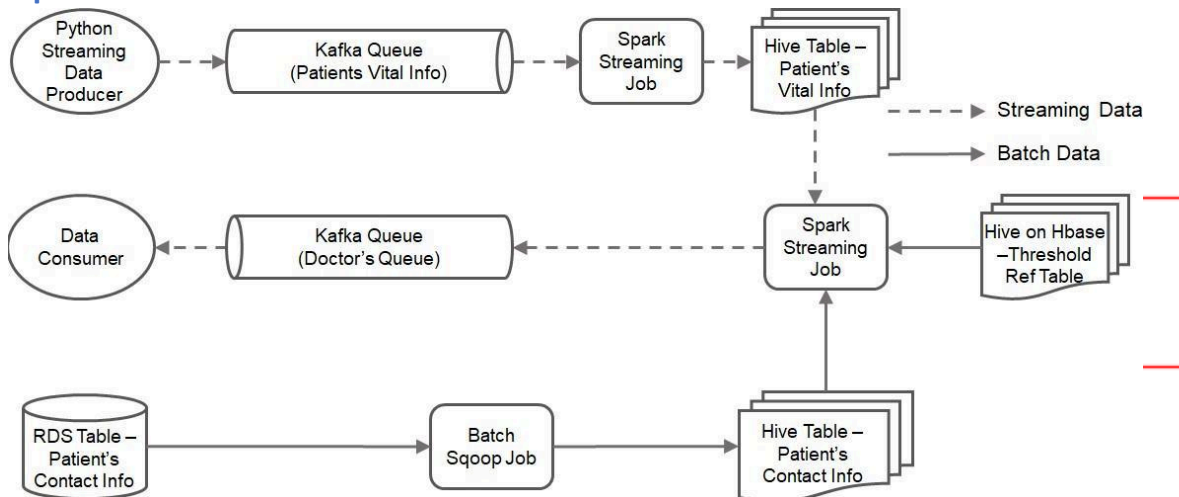
```
select * from Patients_Contact_Info;
```

- View the records in **Patients\_Contact\_Info** table

```
hive> select * from Patients_Contact_Info;
OK
patients_contact_info.patientid patients_contact_info.patientname patients
_contact_info.patientaddress patients_contact_info.phone_number patients
_contact_info.admitted_ward patients_contact_info.age patients_contact
_info.other_details
1 Alex S XDC test Address 8982739282 1 23 null
2 Sammy A New Building Address 2382739282 2 45 null
3 Karan C Aws Address 8923739282 3 56 null
4 Dara M India Address 2182739282 4 67 null
5 Pam ABC test Address 4982739282 5 72 null
Time taken: 1.541 seconds, Fetched: 5 row(s)
```

## PART 2:

Create an HBase table to store threshold reference information and create a hive external table on top of this HBase table



- Create a HBase table named **threshold\_ref** with 3 column families: attribute, limit, alert
- Insert 12 records in this HBase table
- Set up Hive-HBase integration (since HBase and Hive are on separate clusters)
- Create a Hive external table named **Threshold\_Reference\_Table** on top of HBase table

Navigate to HBase shell using below commands:

```
sudo -i
hbase shell
```

```
[root@ip-172-31-92-60 ~]# hbase shell
HBase Shell
Use "help" to get list of supported commands.
Use "exit" to quit this interactive shell.
Version 1.4.13, rUnknown, Fri Apr 17 15:18:24 UTC 2020
```

## Create threshold\_ref table in HBase

```
create 'threshold_ref','attribute','limit','alert'
```

```
hbase(main):001:0> create 'threshold_ref','attribute','limit','alert'  
0 row(s) in 1.6340 seconds
```

### [Insert records into HBase table](#)

```
put 'threshold_ref', '1', 'attribute:attribute', 'heartBeat'  
put 'threshold_ref', '1', 'limit:low_age_limit', '0'  
put 'threshold_ref', '1', 'limit:high_age_limit', '40'  
put 'threshold_ref', '1', 'limit:low_value', '0'  
put 'threshold_ref', '1', 'limit:high_value', '69'  
put 'threshold_ref', '1', 'alert:alert_flag', '1'  
put 'threshold_ref', '1', 'alert:alert_message', 'Low Heart Rate than Normal'  
  
put 'threshold_ref', '2', 'attribute:attribute', 'heartBeat'  
put 'threshold_ref', '2', 'limit:low_age_limit', '0'  
put 'threshold_ref', '2', 'limit:high_age_limit', '40'  
put 'threshold_ref', '2', 'limit:low_value', '70'  
put 'threshold_ref', '2', 'limit:high_value', '78'  
put 'threshold_ref', '2', 'alert:alert_flag', '0'  
put 'threshold_ref', '2', 'alert:alert_message', 'Normal'  
  
put 'threshold_ref', '3', 'attribute:attribute', 'heartBeat'  
put 'threshold_ref', '3', 'limit:low_age_limit', '0'  
put 'threshold_ref', '3', 'limit:high_age_limit', '40'  
put 'threshold_ref', '3', 'limit:low_value', '79'  
put 'threshold_ref', '3', 'limit:high_value', '9999'  
put 'threshold_ref', '3', 'alert:alert_flag', '1'  
put 'threshold_ref', '3', 'alert:alert_message', 'Higher Heart Rate than Normal'  
  
put 'threshold_ref', '4', 'attribute:attribute', 'bp'  
put 'threshold_ref', '4', 'limit:low_age_limit', '0'  
put 'threshold_ref', '4', 'limit:high_age_limit', '40'  
put 'threshold_ref', '4', 'limit:low_value', '0'  
put 'threshold_ref', '4', 'limit:high_value', '160'  
put 'threshold_ref', '4', 'alert:alert_flag', '1'  
put 'threshold_ref', '4', 'alert:alert_message', 'Low BP than Normal'  
  
put 'threshold_ref', '5', 'attribute:attribute', 'bp'  
put 'threshold_ref', '5', 'limit:low_age_limit', '0'  
put 'threshold_ref', '5', 'limit:high_age_limit', '40'  
put 'threshold_ref', '5', 'limit:low_value', '161'  
put 'threshold_ref', '5', 'limit:high_value', '220'  
put 'threshold_ref', '5', 'alert:alert_flag', '0'  
put 'threshold_ref', '5', 'alert:alert_message', 'Normal'  
put 'threshold_ref', '6', 'attribute:attribute', 'bp'  
put 'threshold_ref', '6', 'limit:low_age_limit', '0'  
put 'threshold_ref', '6', 'limit:high_age_limit', '40'  
put 'threshold_ref', '6', 'limit:low_value', '221'  
put 'threshold_ref', '6', 'limit:high_value', '9999'  
put 'threshold_ref', '6', 'alert:alert_flag', '1'  
put 'threshold_ref', '6', 'alert:alert_message', 'Higer BP than Normal'  
  
put 'threshold_ref', '7', 'attribute:attribute', 'heartBeat'  
put 'threshold_ref', '7', 'limit:low_age_limit', '41'  
put 'threshold_ref', '7', 'limit:high_age_limit', '100'
```



```
put 'threshold_ref', '7', 'limit:low_value', '0'
put 'threshold_ref', '7', 'limit:high_value', '65'
put 'threshold_ref', '7', 'alert:alert_flag', '1'
put 'threshold_ref', '7', 'alert:alert_message', 'Low Heart Rate than Normal'

put 'threshold_ref', '8', 'attribute:attribute', 'heartBeat'
put 'threshold_ref', '8', 'limit:low_age_limit', '41'
put 'threshold_ref', '8', 'limit:high_age_limit', '100'
put 'threshold_ref', '8', 'limit:low_value', '66'
put 'threshold_ref', '8', 'limit:high_value', '73'
put 'threshold_ref', '8', 'alert:alert_flag', '0'
put 'threshold_ref', '8', 'alert:alert_message', 'Normal'

put 'threshold_ref', '9', 'attribute:attribute', 'heartBeat'
put 'threshold_ref', '9', 'limit:low_age_limit', '41'
put 'threshold_ref', '9', 'limit:high_age_limit', '100'
put 'threshold_ref', '9', 'limit:low_value', '74'
put 'threshold_ref', '9', 'limit:high_value', '9999'
put 'threshold_ref', '9', 'alert:alert_flag', '1'
put 'threshold_ref', '9', 'alert:alert_message', 'Higher Heart Rate than Normal'

put 'threshold_ref', '10', 'attribute:attribute', 'bp'
put 'threshold_ref', '10', 'limit:low_age_limit', '41'
put 'threshold_ref', '10', 'limit:high_age_limit', '100'
put 'threshold_ref', '10', 'limit:low_value', '0'
put 'threshold_ref', '10', 'limit:high_value', '150'
put 'threshold_ref', '10', 'alert:alert_flag', '1'
put 'threshold_ref', '10', 'alert:alert_message', 'Low BP than Normal'

put 'threshold_ref', '11', 'attribute:attribute', 'bp'
put 'threshold_ref', '11', 'limit:low_age_limit', '41'
put 'threshold_ref', '11', 'limit:high_age_limit', '100'
put 'threshold_ref', '11', 'limit:low_value', '151'
put 'threshold_ref', '11', 'limit:high_value', '180'
put 'threshold_ref', '11', 'alert:alert_flag', '0'
put 'threshold_ref', '11', 'alert:alert_message', 'Normal';

put 'threshold_ref', '12', 'attribute:attribute', 'bp'
put 'threshold_ref', '12', 'limit:low_age_limit', '41'
put 'threshold_ref', '12', 'limit:high_age_limit', '100'
put 'threshold_ref', '12', 'limit:low_value', '181'
put 'threshold_ref', '12', 'limit:high_value', '9999'
put 'threshold_ref', '12', 'alert:alert_flag', '1'
put 'threshold_ref', '12', 'alert:alert_message', 'Higher BP than Normal'
```

## Screenshot for insertion of records

```
hbase(main):002:0> put 'threshold_ref', '1', 'attribute:attribute', 'heartBeat'
ert_flag', '1'
put 'threshold_ref', '1', 'alert:alert_message', 'Low Heart Rate than Normal'
put 'threshold_ref', '2', 'attribute:attribute', 'heartBeat'
put 'threshold_ref', '2', 'limit:low_age_limit', '0'
put 'threshold_ref', '2', 'limit:high_age_limit', '40'
put 'threshold_ref', '2', 'limit:low_value', '70'
put 'threshold_ref', '2', 'limit:high_value', '78'
put 'threshold_ref', '2', 'alert:alert_flag', '0'
put 'threshold_ref', '2', 'alert:alert_message', 'Normal'
put 'threshold_ref', '3', 'attribute:attribute', 'heartBeat'
put 'threshold_ref', '3', 'limit:low_age_limit', '0'
put 'threshold_ref', '3', 'limit:high_age_limit', '40'
put 'threshold_ref', '3', 'limit:low_value', '79'
put 'threshold_ref', '3', 'limit:high_value', '9999'
put 'threshold_ref', '3', 'alert:alert_flag', '1'
put 'threshold_ref', '3', 'alert:alert_message', 'Higher Heart Rate than Normal'
put 'threshold_ref', '4', 'attribute:attribute', 'bp'
put 'threshold_ref', '4', 'limit:low_age_limit', '0'
put 'threshold_ref', '4', 'limit:high_age_limit', '40'
put 'threshold_ref', '4', 'limit:low_value', '0'
put 'threshold_ref', '4', 'limit:high_value', '160'
put 'threshold_ref', '4', 'alert:alert_flag', '1'
put 'threshold_ref', '4', 'alert:alert_message', 'Low BP than Normal'
put 'threshold_ref', '5', 'attribute:attribute', 'bp'
put 'threshold_ref', '5', 'limit:low_age_limit', '0'
put 'threshold_ref', '5', 'limit:high_age_limit', '40'
put 'threshold_ref', '5', 'limit:low_value', '161'
put 'threshold_ref', '5', 'limit:high_value', '220'
put 'threshold_ref', '5', 'alert:alert_flag', '0'
put 'threshold_ref', '5', 'alert:alert_message', 'Normal'
put 'threshold_ref', '6', 'attribute:attribute', 'bp'
put 'threshold_ref', '6', 'limit:low_age_limit', '0'
put 'threshold_ref', '6', 'limit:high_age_limit', '40'
put 'threshold_ref', '6', 'limit:low_value', '221'
put 'threshold_ref', '6', 'limit:high_value', '9999'
put 'threshold_ref', '6', 'alert:alert_flag', '1'
put 'threshold_ref', '6', 'alert:alert_message', 'Higer BP than Normal'
put 'threshold_ref', '7', 'attribute:attribute', 'heartBeat'
put 'threshold_ref', '7', 'limit:low_age_limit', '41'
put 'threshold_ref', '7', 'limit:high_age_limit', '100'
```

## View records in HBase table

```
scan 'threshold_ref'
```

```
hbase(main):098:0> scan 'threshold_ref'
```

## Screenshots of records

ROW	COLUMN+CELL
1	column=alert:alert_flag, timestamp=1679727530966, value=1
1	column=alert:alert_message, timestamp=1679727530971, value=Low Heart Rate than Normal
1	column=attribute:attribute, timestamp=1679727530926, value=heartBeat
1	column=limit:high_age_limit, timestamp=1679727530951, value=40
1	column=limit:high_value, timestamp=1679727530962, value=69
1	column=limit:low_age_limit, timestamp=1679727530947, value=0
1	column=limit:low_value, timestamp=1679727530958, value=0
10	column=alert:alert_flag, timestamp=1679727531306, value=1
10	column=alert:alert_message, timestamp=1679727531310, value=Low BP than Normal
10	column=attribute:attribute, timestamp=1679727531291, value=bp
10	column=limit:high_age_limit, timestamp=1679727531297, value=100
10	column=limit:high_value, timestamp=1679727531303, value=150
10	column=limit:low_age_limit, timestamp=1679727531294, value=41
10	column=limit:low_value, timestamp=1679727531300, value=0
11	column=alert:alert_flag, timestamp=1679727531326, value=0
11	column=alert:alert_message, timestamp=1679727531329, value=Normal
11	column=attribute:attribute, timestamp=1679727531312, value=bp
11	column=limit:high_age_limit, timestamp=1679727531318, value=100
11	column=limit:high_value, timestamp=1679727531323, value=180
11	column=limit:low_age_limit, timestamp=1679727531315, value=41
11	column=limit:low_value, timestamp=1679727531321, value=151

```

8      column=alert:alert_message, timestamp=1679992681690, value
      =Normal
8      column=attribute:attribute, timestamp=1679992681657, value
      =heartBeat
8      column=limit:high_age_limit, timestamp=1679992681668, valu
      e=100
8      column=limit:high_value, timestamp=1679992681679, value=73
8      column=limit:low_age_limit, timestamp=1679992681663, value
      =41
8      column=limit:low_value, timestamp=1679992681674, value=66
9      column=alert:alert_flag, timestamp=1679992681728, value=1
9      column=alert:alert_message, timestamp=1679992681734, value
      =Higher Heart Rate than Normal
9      column=attribute:attribute, timestamp=1679992681696, value
      =heartBeat
9      column=limit:high_age_limit, timestamp=1679992681708, valu
      e=100
9      column=limit:high_value, timestamp=1679992681722, value=99
9      column=limit:low_age_limit, timestamp=1679992681702, value
      =41
9      column=limit:low_value, timestamp=1679992681715, value=74
12 row(s) in 0.3570 seconds

```

### [Threshold Reference Table in Hive](#)

```

CREATE EXTERNAL TABLE Threshold_Reference_Table (
    key int,
    Attribute string,
    low_age_limit int,
    high_age_limit int,
    Low_Range_Value int,
    High_Range_Value int,
    Alert_Flag int,
    Alert_Message string
)
STORED BY 'org.apache.hadoop.hive.hbase.HBaseStorageHandler'
WITH SERDEPROPERTIES (
    'hbase.columns.mapping' = ':key, attribute:attribute, limit:low_age_limit, limit:high_age_limit, limit:low_value,
    limit:high_value, alert:alert_flag, alert:alert_message',
    'hbase.table.name' = 'threshold_ref'
)
TBLPROPERTIES ('hbase.mapred.output.outputtable' = 'threshold_ref');

```

```
hive> CREATE EXTERNAL TABLE Threshold_Reference_Table (
  >   key int,
  >   Attribute string,
  >   low_age_limit int,
  >   high_age_limit int,
  >   Low_Range_Value int,
  >   High_Range_Value int,
  >   Alert_Flag int,
  >   Alert_Message string
  > )
  > STORED BY 'org.apache.hadoop.hive.hbase.HBaseStorageHandler'
  > WITH SERDEPROPERTIES (
  >   'hbase.columns.mapping' = ':key, attribute:attribute, limit:low_age_limit, limit:high_age_limit, limit:low_value, limit:high_value, alert:alert_flag, alert:alert_message',
  >   'hbase.table.name' = 'threshold_ref'
  > )
  > TBLPROPERTIES ('hbase.mapred.output.outputtable' = 'threshold_ref');
OK
Time taken: 2.31 seconds
```

- View the contents of Threshold\_Reference\_Table

```
set hive.cli.print.header = true;

SELECT * FROM Threshold_Reference_Table order by key;
```

```
hive> set hive.cli.print.header = true;
hive> select * from Threshold_Reference_Table order by key;
Query ID = hadoop_20230328084446_4e04390b-9e15-4799-a8f7-ef03baf510dd
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1679992808558_0001)

Map 1: -/-      Reducer 2: 0/1
Map 1: 0/1      Reducer 2: 0/1
Map 1: 0/1      Reducer 2: 0/1
Map 1: 0(+1)/1  Reducer 2: 0/1
Map 1: 1/1      Reducer 2: 0(+1)/1
Map 1: 1/1      Reducer 2: 1/1
```

Screenshot of Threshold\_Reference\_Table records:

```
OK
threshold_reference_table.key  threshold_reference_table.attribute  threshol
d_reference_table.low_age_limit threshold_reference_table.high_age_limit  t
hreshold_reference_table.low_range_value  threshold_reference_table.high_r
ange_value  threshold_reference_table.alert_flag  threshold_reference_tabl
e.alert_message
1      heartBeat      0      40      0      69      1      Low Heart Rate t
han Normal
2      heartBeat      0      40      70      78      0      Normal
3      heartBeat      0      40      79      9999      1      Higher Heart Rat
e than Normal
4      bp      0      40      0      160      1      Low BP than Normal
5      bp      0      40      161      220      0      Normal
6      bp      0      40      221      9999      1      Higer BP than Normal
7      heartBeat      41      100      0      65      1      Low Heart Rate t
han Normal
8      heartBeat      41      100      66      73      0      Normal
9      heartBeat      41      100      74      9999      1      Higher Heart Rat
e than Normal
10     bp      41      100      0      150      1      Low BP than Normal
11     bp      41      100      151      180      0      Normal
12     bp      41      100      181      9999      1      Higher BP than Normal
Time taken: 15.749 seconds, Fetched: 12 row(s)
```

- Create copy of threshold table *Threshold\_Reference* in Hive and insert the records from Threshold\_Reference\_Table to Threshold\_Reference

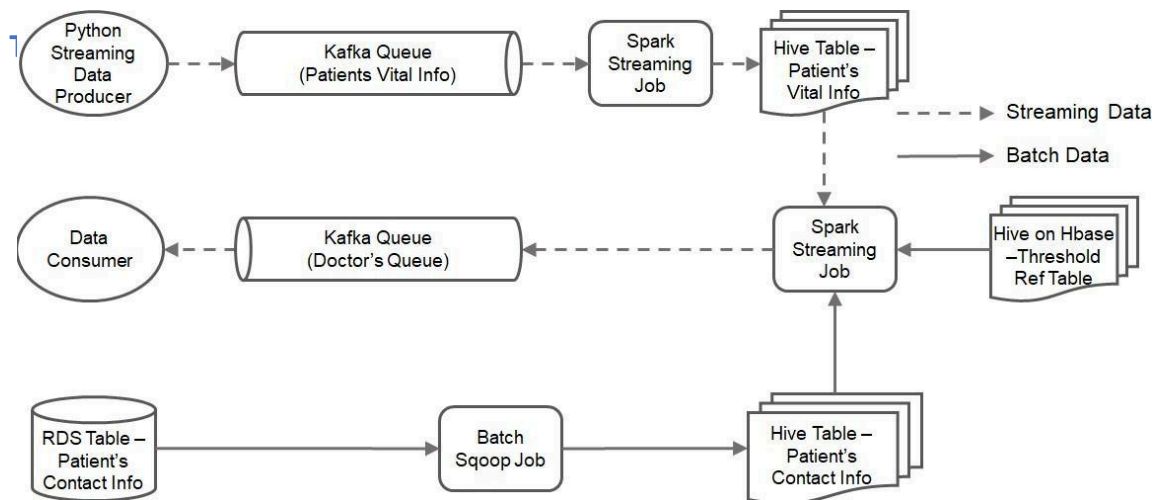
This table can be accessed by spark streaming application 2 for Part 4 mentioned below)

```
CREATE EXTERNAL TABLE Threshold_Reference (
  key int,
  Attribute string,
  low_age_limit int,
  high_age_limit int,
  Low_Range_Value int,
  High_Range_Value int,
  Alert_Flag int,
  Alert_Message string
);

INSERT INTO table Threshold_Reference SELECT * FROM Threshold_Reference_Table;
```

### PART 3:

#### Dealing with Kafka part in EMR Cluster



This involves following tasks:

1. Build a Kafka Producer application in Python to simulate streaming data by reading from RDS every second.
2. Push patient vitals data to the Kafka topic `patients\_vital\_info`.
3. Create a Spark streaming job to:
4. Read data from the Kafka topic.
5. Add a timestamp column.
6. Store the data in HDFS in Parquet format.
7. Set up an external Hive table, `patients\_vital\_info`, to read streaming data from the HDFS location.
8. The producer will process all 1800 records within 30 minutes

## Setting Up Apache Kafka on an EMR Cluster

### Downloading Kafka files

```
wget https://downloads.apache.org/kafka/3.6.2/kafka_2.12-3.6.2.tgz
```

```
hadoop@ip-172-31-5-58:~$ wget https://downloads.apache.org/kafka/3.6.2/kafka_2.12-3.6.2.tgz
--2024-11-05 07:15:56-- https://downloads.apache.org/kafka/3.6.2/kafka_2.12-3.6.2.tgz
Resolving downloads.apache.org (downloads.apache.org)... 88.99.208.237, 135.181.214.104, 2a01:4f9:3a:2c57::2, ...
Connecting to downloads.apache.org (downloads.apache.org)[88.99.208.237]:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 113986658 (109M) [application/x-gzip]
Saving to: 'kafka_2.12-3.6.2.tgz'

64% [=====] 73,135,871 22.2MB/s eta 2s
```

### Extract Kafka files

```
tar -xvf kafka_2.12-3.6.2.tgz
```

```
hadoop@ip-172-31-5-58:~$ tar -xvf kafka_2.12-3.6.2.tgz
kafka_2.12-3.6.2/libs/jetty-client-9.4.53.v20231009.jar
kafka_2.12-3.6.2/libs/maven-artifact-3.8.8.jar
kafka_2.12-3.6.2/libs/swagger-annotations-2.2.8.jar
kafka_2.12-3.6.2/libs/jersey-container-servlet-core-2.39.1.jar
kafka_2.12-3.6.2/libs/jersey-server-2.39.1.jar
kafka_2.12-3.6.2/libs/jersey-client-2.39.1.jar
kafka_2.12-3.6.2/libs/jersey-common-2.39.1.jar
kafka_2.12-3.6.2/libs/jakarta.ws.rs-api-2.1.6.jar
kafka_2.12-3.6.2/libs/hk2-locator-2.6.1.jar
kafka_2.12-3.6.2/libs/javassist-2.29.2-GA.jar
kafka_2.12-3.6.2/libs/javax.activation-api-1.2.0.jar
kafka_2.12-3.6.2/libs/javax.servlet-api-3.1.0.jar
kafka_2.12-3.6.2/libs/jetty-http-9.4.53.v20231009.jar
kafka_2.12-3.6.2/libs/jetty-io-9.4.53.v20231009.jar
kafka_2.12-3.6.2/libs/jetty-util-ajax-9.4.53.v20231009.jar
kafka_2.12-3.6.2/libs/jetty-continuation-9.4.53.v20231009.jar
kafka_2.12-3.6.2/libs/jetty-util-9.4.53.v20231009.jar
kafka_2.12-3.6.2/libs/jsr305-3.0.2.jar
kafka_2.12-3.6.2/libs/plexus-utils-3.3.1.jar
kafka_2.12-3.6.2/libs/commons-lang3-3.8.1.jar
kafka_2.12-3.6.2/libs/hk2-api-2.6.1.jar
kafka_2.12-3.6.2/libs/hk2-utils-2.6.1.jar
kafka_2.12-3.6.2/libs/jakarta.inject-2.6.1.jar
kafka_2.12-3.6.2/libs/jakarta.annotation-api-1.3.5.jar
kafka_2.12-3.6.2/libs/osgi-resource-locator-1.0.3.jar
kafka_2.12-3.6.2/libs/jakarta.validation-api-2.0.2.jar
kafka_2.12-3.6.2/libs/aopalliance-repackaged-2.6.1.jar
kafka_2.12-3.6.2/libs/trogdor-3.6.2.jar
kafka_2.12-3.6.2/libs/kafka-shell-3.6.2.jar
kafka_2.12-3.6.2/libs/jline-3.25.1.jar
kafka_2.12-3.6.2/libs/connect-file-3.6.2.jar
kafka_2.12-3.6.2/libs/connect-basic-auth-extension-3.6.2.jar
kafka_2.12-3.6.2/libs/javax.annotation-api-1.3.2.jar
kafka_2.12-3.6.2/libs/connect-mirror-3.6.2.jar
kafka_2.12-3.6.2/libs/connect-mirror-client-3.6.2.jar
kafka_2.12-3.6.2/libs/kafka-streams-3.6.2.jar
kafka_2.12-3.6.2/libs/rocksdbjni-7.9.2.jar
kafka_2.12-3.6.2/libs/kafka-streams-scala-2.12-3.6.2.jar
kafka_2.12-3.6.2/libs/kafka-streams-test-utils-3.6.2.jar
kafka_2.12-3.6.2/libs/kafka-streams-examples-3.6.2.jar
[hadoop@ip-172-31-5-58 ~]$
```

Update the advertised listeners to use our EMR IP address



```
hadoop@ip-172-31-5-58:~/ka x + v
# (the "License"); you may not use this file except in compliance with
# the License. You may obtain a copy of the license at
#
# http://www.apache.org/licenses/LICENSE-2.0
#
# Unless required by applicable law or agreed to in writing, software
# distributed under the license is distributed on an "AS IS" BASIS,
# WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
# See the License for the specific language governing permissions and
# limitations under the License.
#
# This configuration file is intended for use in ZK-based mode, where Apache ZooKeeper is required.
# See kafka.server.KafkaConfig for additional details and defaults
#
##### Server Basics #####
# The id of the broker. This must be set to a unique integer for each broker.
broker.id=0
##### Socket Server Settings #####
# The address the socket server listens on. If not configured, the host name will be equal to the value of
# java.net.InetAddress.getCanonicalHostName(), with PLAINTEXT listener name, and port 9092.
#
# FORMAT:
#   listeners = listener_name://host_name:port
#
# EXAMPLE:
#   listeners = PLAINTEXT://your.host.name:9092
#listeners=PLAINTEXT://:9092
#
# Listener name, hostname and port the broker will advertise to clients.
# If not set, it uses the value for "listeners".
advertised.listeners=PLAINTEXT://ec2-3-239-100-180.compute-1.amazonaws.com:9092
#
# Maps listener names to security protocols, the default is for them to be the same. See the config documentation for more details
#listener.security.protocol.map=PLAINTEXT:PLAINTEXT,SSL:SSL,SASL_PLAINTEXT:SASL_PLAINTEXT,SASL_SSL:SASL_SSL
#
# The number of threads that the server uses for receiving requests from the network and sending responses to the network
num.network.threads=3
```

- Start Zookeeper server

```
cd kafka_2.12-3.6.2
bin/zookeeper-server-start.sh config/zookeeper.properties
```

- Start Kafka server

```
cd downloads/kafka_2.12-3.6.2
bin/kafka-server-start.sh config/server.properties
```

```
cd kafka_2.12-3.6.2
bin/kafka-topics.sh --bootstrap-server localhost:9092 --delete --topic patients_vital_info
```

- Delete the topic if it already exists and create again

```
bin/kafka-topics.sh --create --bootstrap-server localhost:9092 --replication-factor 1 --partitions 1
--topic patients_vital_info
```

- Transfer the Python producer file and Install required packages

```
WINSXP transfer file to emr cluster
pip install mysql-connector-python kafka-python
```

- Run the producer script

```
python kafka_produce_patient_vitals.py
```

## Spark Streaming Job 1 (kafka\_spark\_patient\_vitals.py)

It reads data from Kafka topic and add the **timestamp column** and store in HDFS location in parquet format.

```
export SPARK_KAFKA_VERSION=0.10
spark-submit --packages org.apache.spark:spark-sql-kafka-0-10_2.11:2.4.5 kafka_spark_patient_vitals.py
```

```
hive

use health;

CREATE EXTERNAL TABLE health.patients_vital_info (
  customerId int,
  heartBeat int,
  bp int,
  message_time timestamp)
STORED AS PARQUET
LOCATION '/user/hadoop/health-alert/patients-vital-info/'
TBLPROPERTIES ('parquet.compress'='SNAPPY');
```

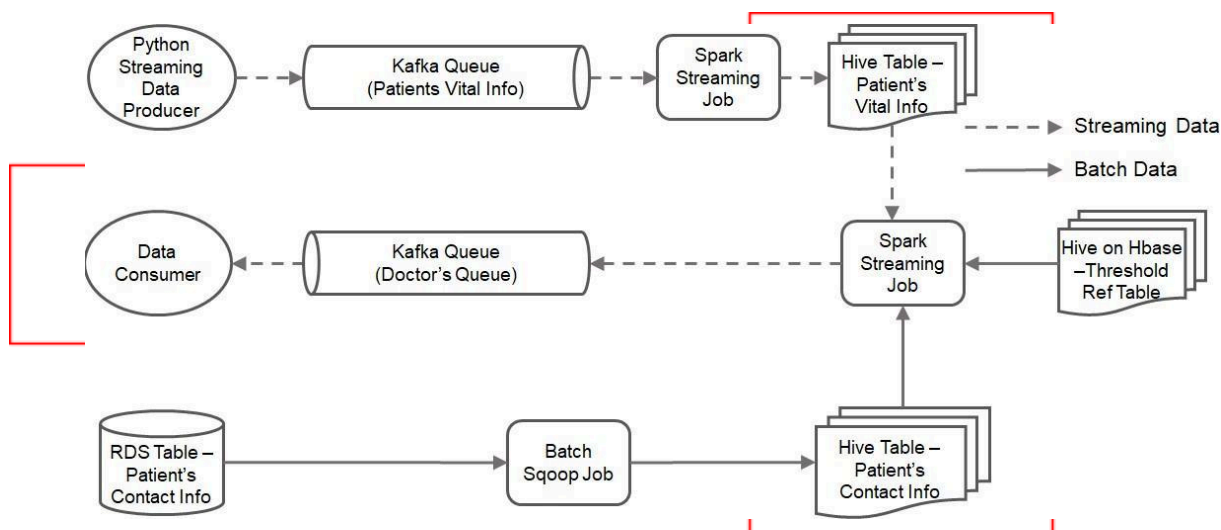
## In the Hive Shell

- Create external table 'patients\_vital\_info'

```
hive> CREATE EXTERNAL TABLE health.patients_vital_info (
>   customerId int,
>   heartBeat int,
>   bp int,
>   message_time timestamp)
> STORED AS PARQUET
> LOCATION '/user/hadoop/health-alert/patients-vital-info/'
> TBLPROPERTIES ('parquet.compress'='SNAPPY');
OK
Time taken: 0.055 seconds
```

## PART 4:

Compare vital information with threshold values, analyze the results, and send notifications if the data exceeds threshold limits.



This involves following tasks:

1. Create a Kafka topic called `alerts_message` to store abnormal patient vital signs.
2. Develop a Spark streaming job to:
3. Read data from three Hive tables.
4. Analyze patient vitals and, if irregular, push them to the `alerts_message` Kafka topic.
5. Set up a Kafka consumer to read messages from the `alerts_message` topic.
6. Use SNS to send email notifications for messages received by the consumer.

## Create a Kafka topic ( alerts\_message)

- Delete the topic if it already exists

```
bin/kafka-topics.sh --bootstrap-server localhost:9092 --delete --topic alerts_message
```

- Create the topic again

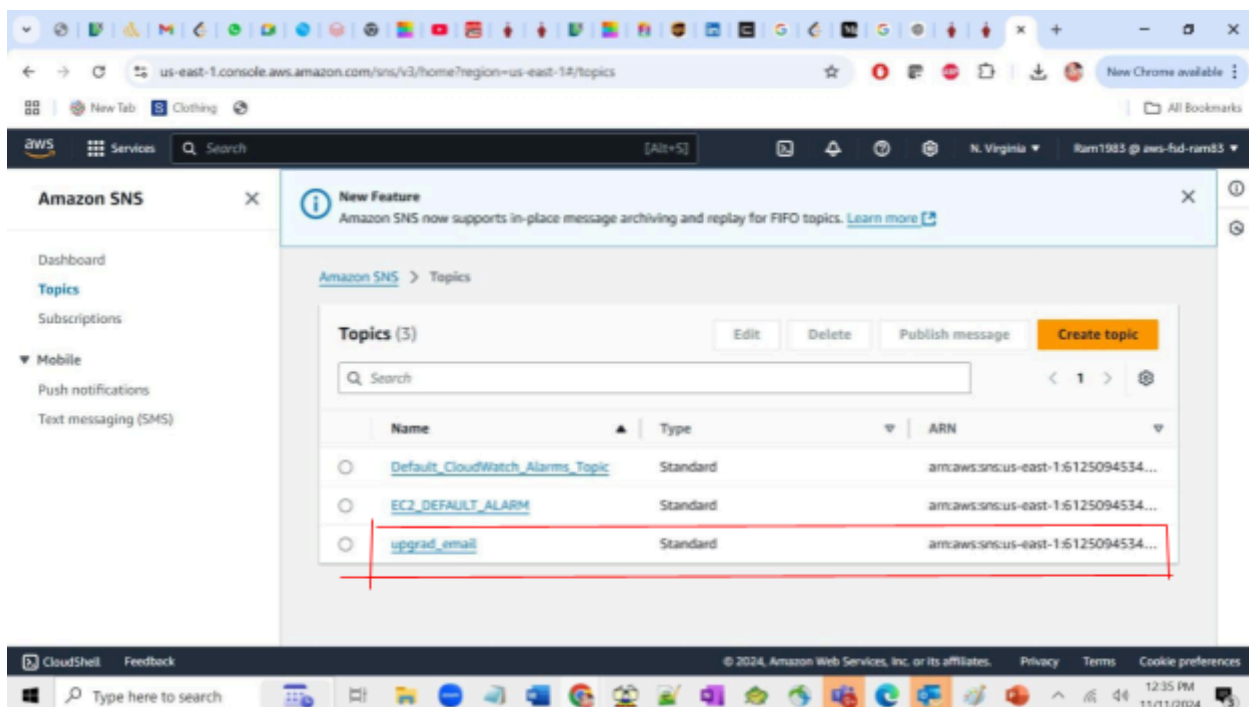
```
bin/kafka-topics.sh --create --bootstrap-server localhost:9092 --replication-factor 1 --partitions 1 --topic alerts_message
```

## Spark Streaming Job to push irregular patient vitals to alerts\_message Kafka topic

```
export SPARK_KAFKA_VERSION=0.10
spark-submit --executor-memory 4G --num-executors 4 --packages
org.apache.spark:spark-sql-kafka-0-10_2.11:2.4.5
kafka_spark_generate_alerts.py
```

## Configure SNS

### Create SNS topic (Health\_Alerts)



## Subscribe to SNS topic

The screenshot shows the Amazon SNS console interface. On the left is a navigation sidebar with links to Dashboard, Topics, Subscriptions, Mobile, Push notifications, and Text messaging (SMS). The main content area is titled "Subscription: 025aac40-c3ed-4abe-9fdf-1db154806721" and includes "Edit" and "Delete" buttons. Below this is a "Details" section with the following information:

ARN arn:aws:sns:us-east-1:612509453432:upgrad_email:025aac40-c3ed-4abe-9fdf-1db154806721	Status Confirmed
Endpoint rama.r.krishnan@gmail.com	Protocol EMAIL
Topic upgrad_email	
Subscription Principal	

## Subscription Confirmation Email

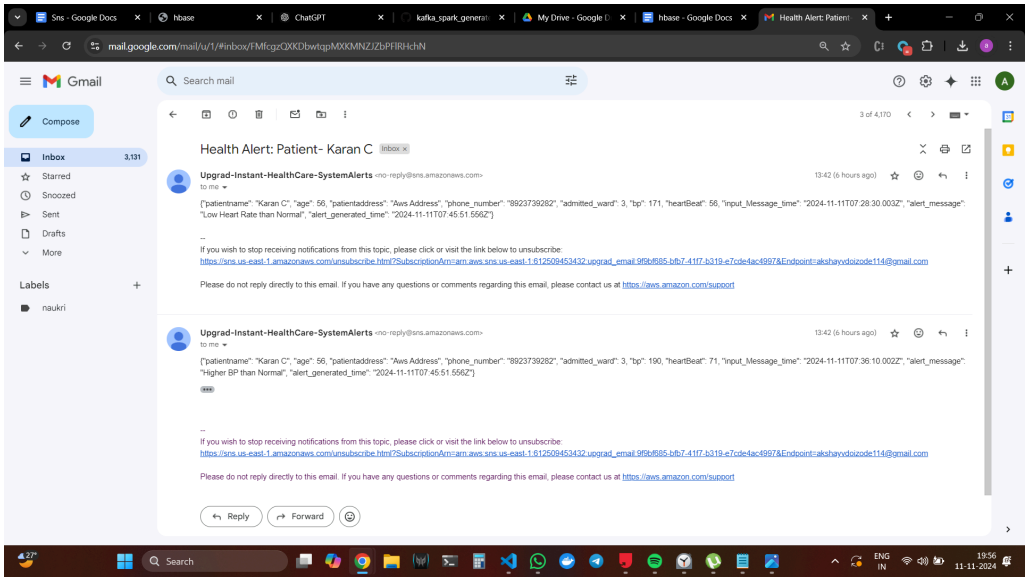
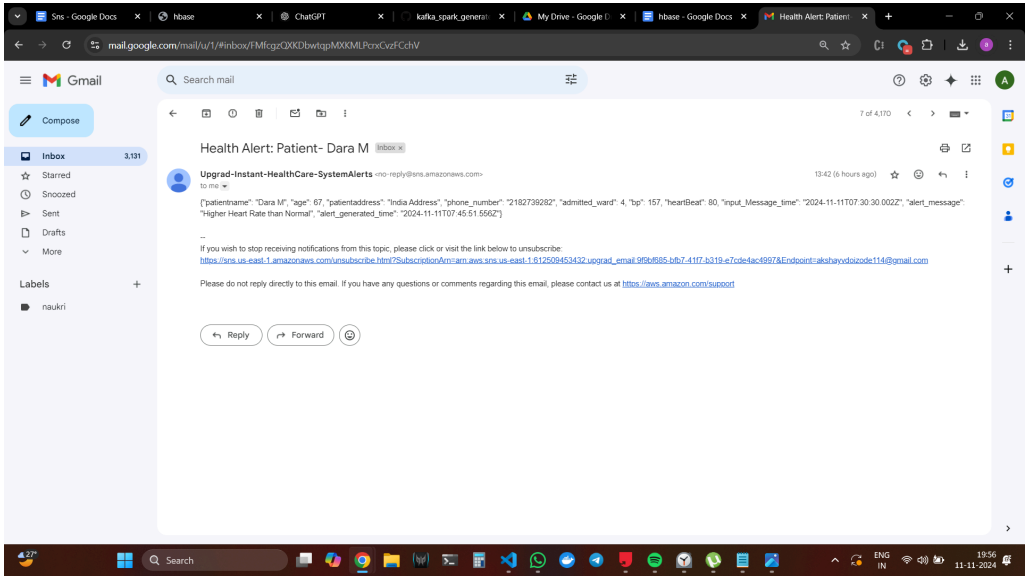
The screenshot shows an email titled "AWS Notification - Subscription Confirmation". The sender is "AWS Notifications" with the email address "no-reply@sns.amazonaws.com". The email content states: "You have chosen to subscribe to the topic: arn:aws:sns:us-east-1:612509453432:upgrad\_email". It includes a link to "Confirm subscription" and a note: "To confirm this subscription, click or visit the link below (if this was in error no action is necessary)." At the bottom, there are buttons for "Reply", "Forward", and a smiley face icon.

## Kafka consumer

```
pip install kafka-python
pip install boto3
python kafka_consume_alerts.py
```

## Final Output Screenshot:

### Patient health Alert email - using SNS



## Summary

### 1. Setup and Configuration:

- **Kafka Installation:** Download and extract Kafka, configure the server, and start both Zookeeper and Kafka servers. Create two Kafka topics: `patients\_vital\_info` (for patient data) and `alerts\_message` (for alerts).
- **Python and Kafka Libraries:** Install required Python libraries (`mysql-connector-python` and `kafka-python`) for data processing and connectivity.
- **MySQL Setup:** Install MySQL Connector, configure MySQL, set up user privileges, and enable access for data import.

### 2. Data Import with Sqoop and Storage in Hive:

- **MySQL Data Import:** Import patient contact information from MySQL RDS into HDFS using Sqoop.
- **Hive Table Creation:** Create a Hive database and external table (`patients\_contact\_info`) to store imported patient data

### **3. Threshold Configuration with HBase and Hive Integration:**

- HBase Table Creation: Create `threshold\_ref` in HBase to store threshold values for patient vitals.
- Hive-HBase Integration: Create an external Hive table (`threshold\_reference\_table`) linked to `threshold\_ref` in HBase, allowing Hive to access threshold data for analysis.

### **4. Streaming and Real-Time Processing with Kafka and Spark:**

- Kafka Producer and Spark Job: A Python Kafka producer simulates streaming data by sending patient vitals to `patients\_vital\_info`. A Spark streaming job then reads data from Kafka, adds timestamps, and stores it in HDFS in Parquet format.
- Hive Table for Streaming Data: Create an external Hive table (`patients\_vital\_info`) to read the streaming data from HDFS.

### **5. Alert Generation and Notification:**

- Spark Streaming for Alerts: Create a Spark job to analyze the patient vitals against threshold values from Hive tables. Any irregularities are pushed to the `alerts\_message` Kafka topic.
- Kafka Consumer and SNS Notifications: Set up a Kafka consumer to read from `alerts\_message` and send email notifications using AWS SNS for any alert messages.